

Truss Decomposition on Shared-Memory Parallel Systems

Shaden Smith^{1,2}, Xing Liu², Nesreen K. Ahmed²,
Ancy Sarah Tom¹, Fabrizio Petrini², and George Karypis¹

¹Department of Computer Science & Engineering, University of Minnesota


²Intel Parallel Computing Lab

shaden@cs.umn.edu

GraphChallenge Finalist, HPEC 2017

Truss decomposition

We are interested in computing the complete *truss decomposition* of a graph on shared-memory parallel systems.



Notation:

- ▶ A k -truss is a subgraph in which each edge is contained in at least $(k-2)$ triangles in the same subgraph.
- ▶ The truss number of an edge, $\Gamma(e)$, is the maximum k -truss that contains e .

Serial peeling algorithm

Peeling builds the truss decomposition bottom-up.

```
1: Compute initial supports and store in  $sup(\cdot)$ 
2:  $k \leftarrow 3$ 
3: while  $|E| > 0$  do
4:   for each edge  $e$  not in current  $k$ -truss do
5:     for each edge  $e' \in \Delta_e$  do
6:        $sup(e') \leftarrow sup(e') - 1$ 
7:     end for
8:      $\Gamma(e) \leftarrow k - 1$ 
9:     Remove  $e$  from  $E$ 
10:  end for
11:   $k \leftarrow k + 1$ 
12: end while
```

Multi-Stage Peeling (MSP)


We break the peeling process into several bulk-synchronous substeps.

High-level idea:

- ▶ Store the graph as an adjacency list for each vertex (i.e., CSR).
- ▶ Do a 1D decomposition on the vertices.
- ▶ Operations which modify graph state (e.g., edge deletion and support updates) are grouped by source vertex.
 - ▶ Batching localizes updates to a specific adjacency list and eliminates race conditions.


Multi-Stage Peeling (MSP)

Step 1: frontier generation




Multi-Stage Peeling (MSP)

Step 2: triangle enumeration




Multi-Stage Peeling (MSP)

Step 3: support updates



Multi-Stage Peeling (MSP)

Step 4: edge deletion



Experimental Setup

Software:

- ▶ Parallel baseline: asynchronous nucleus decomposition (AND)¹, written in C and parallelized with OpenMP
- ▶ MSP is written in C and parallelized with OpenMP
- ▶ Compiled with `icc v17.0`

Hardware:

- ▶ 56-core shared-memory system ($2 \times$ 28-core Skylake Xeon)
- ▶ 192GB DDR4 memory

¹A. E. Sariyuce, C. Seshadhri, and A. Pinar, “*Parallel local algorithms for core, truss, and nucleus decompositions*,” arXiv preprint arXiv:1704.00386, 2017.

Graphs

More datasets in paper.

Graph	$ V $	$ E $	$ \Delta $	k_{\max}
cit-Patents	3.8M	16.5M	7.5M	36
soc-Orkut	3.0M	106.3M	524.6M	75
twitter	41.7M	1.2B	34.8B	1998
rmat22	2.4M	64.1M	2.1B	485
rmat23	4.5M	129.3M	4.5B	625
rmat24	8.9M	260.3M	9.9B	791
rmat25	17.0M	523.5M	21.6B	996

K, **M**, and **B** denote thousands, millions, and billions, respectively. The first group of graphs is taken from real-world datasets, and the second group is synthetic.

Parallel baseline comparison

MSP is up to $28\times$ faster than AND and $20\times$ faster than the serial peeling algorithm.

Graph	Peeling	AND		MSP	
cit-Patents	2.89	0.23	$12.6\times$	0.58	$5.0\times$
soc-Orkut	228.06	64.31	$3.5\times$	11.30	$20.2\times$
twitter	-	-	-	1566.72	
rmat22	403.59	398.46	$1.0\times$	42.22	$9.6\times$
rmat23	980.68	1083.66	$0.9\times$	85.14	$11.5\times$
rmat24	2370.54	4945.70	$0.5\times$	175.29	$13.5\times$
rmat25	5580.47	-	-	352.37	$15.8\times$

Values are runtimes, in seconds, of the full truss decomposition. **Peeling** is the optimized serial implementation. **AND** and **MSP** are executed on 56 cores.

Wrapping up

Multi-stage peeling (MSP):

- ▶ processes graph mutations in batches to avoid race conditions
 - ▶ resulting algorithm is free of atomics and mutexes
- ▶ can decompose a billion-scale graph on a single node in minutes

Relative to the state-of-the-art:

- ▶ Up to $28\times$ speedup over the state-of-the-art parallel algorithm
- ▶ Serial optimizations achieve over $1400\times$ speedup over the provided Matlab benchmark (*in paper*).

shaden@cs.umn.edu

Backup

Peeling algorithm

```
1: Compute initial supports and store in  $sup$ 
2:  $k \leftarrow 3$ 
3: while  $|E| > 0$  do
4:    $\mathcal{F}_k \leftarrow \{e \in E : sup(e) < k - 2\}$ 
5:   while  $|\mathcal{F}_k| > 0$  do
6:     for  $e \in \mathcal{F}_k$  do
7:       for  $e' \in \Delta_e$  do
8:          $sup(e') \leftarrow sup(e') - 1$ 
9:       end for
10:       $E \leftarrow E \setminus \{e\}$ 
11:       $\Gamma(e) \leftarrow k - 1$ 
12:       $\mathcal{F}_k \leftarrow \{e \in E : sup(e) < k - 2\}$ 
13:    end for
14:  end while
15:   $k \leftarrow k + 1$ 
16: end while
```

Parallelization challenges

A natural first approach to parallelization is to peel edges concurrently.

There are several challenges when parallelizing:

- ▶ graph data structure is dynamic
- ▶ supports must be decremented safely
- ▶ triangles may be counted multiple times


Serial benchmark comparison

The optimized peeling implementation achieves $1400\times$ speedup over the GraphChallenge benchmark (both serial).


Graph	Octave	Peeling	Speedup
soc-Slashdot0811	169.23	0.22	$769.1\times$
cit-HepTh	448.23	0.40	$1120.6\times$
soc-Epinions1	675.03	0.46	$1467.4\times$
loc-gowalla	787.95	0.79	$997.4\times$
cit-Patents	972.66	4.03	$241.4\times$

Values are runtime in seconds. **Octave** is the serial Octave benchmark provided by the GraphChallenge specification. **Peeling** is the proposed serial implementation of the peeling algorithm. Speedup is measured relative to **Octave**.


Serial breakdown



Parallel breakdown




Strong scaling



Cost per truss

The time per k -truss on soc-orkut is unsurprising.



Cost per truss

rmat25 is more challenging.

